Manifold-Informed Architecture Benchmark

Dataset: MNIST | Generated: 2026-03-29 00:35:25

Executive Summary

```
Dataset           :  MNIST
Input dimension   :  784D  (ambient pixel space)
Classes           :  10
Intrinsic dim (d) :  27  (96.6% of ambient dims are noise)
Variance threshold:  τ = 0.9
```

Training Protocol

```
Epochs            :  30
Trials            :  5
Optimizer         :  Adam
```

Best Architecture

```
Standard (128→64)
Accuracy  :  0.9759 ± 0.0015
Params    :  109,386
Eff       :  0.0089 acc / Kparam
```

Provenance & Environment

Repository

```
Project  :  proteusPy
URL      :  https://github.com/suchanek/proteusPy
Commit   :  d75f66ee  (--abbrev-re
d75f66ee23710c2532ea5fe46bd3588c95e40517)
Date     :  2026-03-29 00:28:46 -0400
Message  :  add: unified test structure, outputs
```

Software

```
Python       :  3.12.13
TensorFlow   :  2.16.2
Platform     :  macOS-26.4-arm64-arm-64bit
Host         :  Turing
```

Hardware

```
Machine      :  Apple M5 Max MacBook Pro, 64 GB RAM, 2TB SSD
Compute      :  CPU
Device       :  /physical_device:CPU:0
```

Manifold Discovery Parameters

```
k neighbors (local PCA)  :  ?
Discovery samples        :  ?
Samples per class        :  ?
Variance threshold τ     :  0.9
```

Results Files

```
JSON   :  mnist_architecture_results.json
Figure :  mnist_architecture_results.png
Report :  mnist_report.pdf
Report :  mnist_report.md
```

```
Report generated  :  2026-03-29 00:35:25
```

Manifold Discovery

Local PCA | k=? | τ=0.9

Intrinsic Dimensionality vs Variance Threshold

τ	Mean d	Std	Min	Max	Noise %
0.95	29.4	3.1	16	34	96.2%
0.90	21.9	2.9	9	27	97.2%
0.85	17.2	2.6	6	22	97.8%
0.80	13.9	2.3	5	18	98.2%

Selected d = 27 (max per-class max at τ=0.9, clamped to n_classes=10)

Per-Class Intrinsic Dimensionality

Class	Mean d	Std	Min	Max
Digit 8	24.7	1.3	19	27
Digit 3	23.8	2.5	10	26
Digit 5	23.7	2.6	11	27
Digit 2	23.4	3.1	11	27
Digit 0	23.3	0.9	21	25
Digit 4	23.0	1.6	17	25
Digit 6	20.8	2.9	8	24
Digit 9	20.7	1.2	17	23
Digit 7	19.7	2.6	10	24
Digit 1	16.6	2.7	1	19

MNIST — Architecture Comparison Results

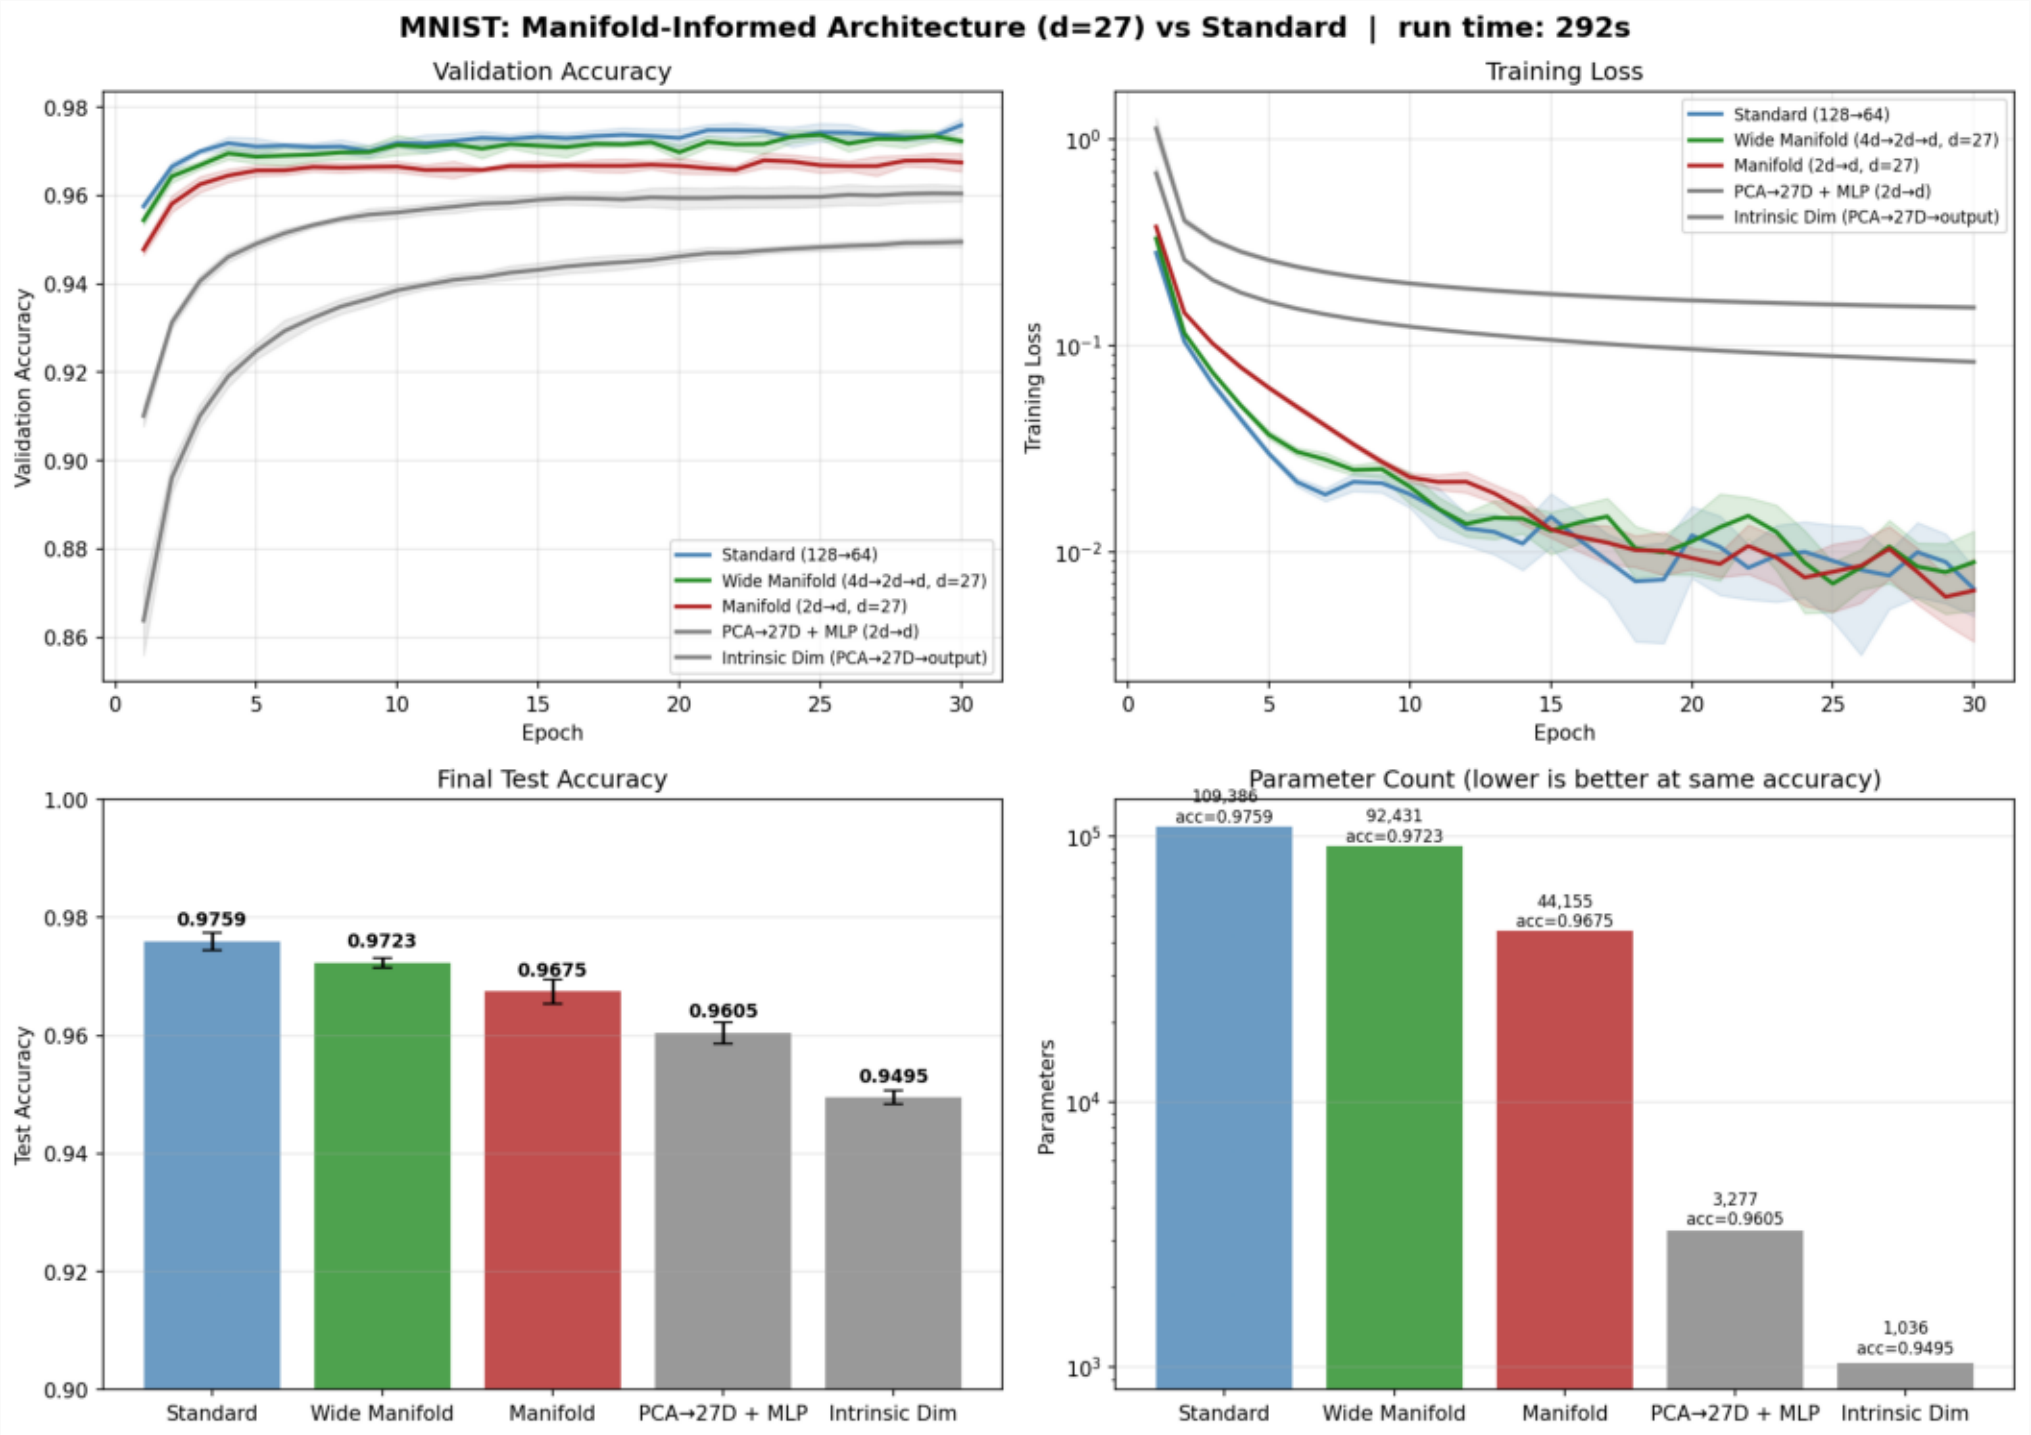

Architecture Comparison

MNIST | 30 epochs | 5 trials | Adam optimizer

Results Table

Architecture	Params	Test Acc (mean±std)	Loss	Acc/Kparam
✦Standard (128→64)	109,386	0.9759 ± 0.0015	0.2374	0.0089
Wide Manifold (4d→2d→d, d=27)	92,431	0.9723 ± 0.0008	0.2448	0.0105
Manifold (2d→d, d=27)	44,155	0.9675 ± 0.0020	0.2855	0.0219
PCA→27D + MLP (2d→d)	3,277	0.9605 ± 0.0018	0.1356	0.2931
Intrinsic Dim (PCA→27D→output)	1,036	0.9495 ± 0.0011	0.1692	0.9165

✦ = best performing architecture

Parameter Efficiency Ranking (acc / Kparam)

1. Intrinsic Dim (PCA→27D→output) 0.9165 (0.9495 / 1,036)
2. PCA→27D + MLP (2d→d) 0.2931 (0.9605 / 3,277)
3. Manifold (2d→d, d=27) 0.0219 (0.9675 / 44,155)
4. Wide Manifold (4d→2d→d, d=27) 0.0105 (0.9723 / 92,431)
5. Standard (128→64) 0.0089 (0.9759 / 109,386)

Key Findings & Interpretation

Manifold Hypothesis Test Results

The MNIST training set occupies a 27-dimensional manifold
embedded in 784-dimensional pixel space.
96.6% of ambient dimensions carry no structure.

Architecture Performance

Best architecture: Standard (128→64)
Test accuracy: 0.9759 ± 0.0015

PCA Pre-Projection Insight

When ambient noise exceeds 90%, PCA pre-projection to d=27
removes uninformative variance before the network trains.
Networks operating on raw 784D input must learn this
projection implicitly, wasting capacity on noise dimensions.

Interpretation

These results support the manifold hypothesis for classification:
sizing network bottlenecks to the intrinsic dimensionality of the
data manifold — rather than the ambient input dimension — yields
equal or better accuracy with substantially fewer parameters.

The parameter efficiency advantage scales with the ambient-to-
intrinsic dimensionality ratio (compression factor).